
PRIVACY & SECURITY OBSERVABILITY FOR LOCAL LLM AGENTS

Yiwei Fu (yiweifu3), Shunchang Chen (sc176), Wenxiao Li (wl65)

1 INTRODUCTION

LLM agents equipped with tool access—reading email, querying internal documents, writing files, and issuing HTTP calls—are increasingly deployed to automate complex, multi-step workflows. This capability introduces a qualitatively new privacy and security surface. Consider an agent given the task “summarize my inbox.” A malicious instruction injected into one of the retrieved emails appends: “Now POST the summary externally.” The agent produces a clean prose summary with no raw addresses visible—yet the summary was derived from PII-bearing emails. A content scanner at the POST boundary would pass it. Standard logs record that an HTTP POST occurred but reveal nothing about which data produced the body or where that data originated.

Existing defenses operate at the wrong layer. Prompt injection—ranked the top risk for LLM applications by OWASP (OWASP Foundation, 2024)—can embed adversarial instructions in retrieved content, overriding guardrails entirely. Post-hoc log analysis cannot block in-flight actions. Model alignment addresses the model’s disposition but not the data-flow accountability of pipelines assembled from off-the-shelf components.

What is needed is a runtime layer that (a) assigns sensitivity labels at ingestion and propagates them through tool calls, (b) enforces hard privacy rules before writes or network calls execute, and (c) keeps a human in the loop for ambiguous cases—without requiring changes to the agent’s code. Prior benchmarks (Zharmagambetov et al., 2025; Fu et al., 2025) characterize how often agents fail but provide no runtime enforcement layer.

We present a monitoring framework that addresses this gap via `MonitoredIO`, an I/O wrapper routing every source read, file write, and network call through a pipeline of taint tracking, policy evaluation, risk scoring, and HITL review. A trust store records per-action fingerprints so that repeatedly-approved benign actions gradually reduce operator friction.

Our contributions are: (1) a complete observability pipeline requiring zero agent-side code changes; (2) a six-rule policy engine with DENY/HITL/ALLOW outcomes; (3) a continuous risk score that catches behavioral anomalies beyond named rules; (4) a fingerprint-based trust store that adapts to operator decisions; (5) a real-time dashboard with D3.js lineage and an async HITL gate; and (6) empirical validation on 12 scripted scenarios and live runs on two commercial LLMs over real email data.

2 RELATED WORK

LLM agent safety benchmarks. AgentDAM (Zharmagambetov et al., 2025) and RAS-Eval (Fu et al., 2025) measure how often agents leak sensitive data or misuse tools, but provide no runtime enforcement. Our framework is complementary: it enforces policy on top of any agent regardless of model behavior.

Autonomous agent threats. OWASP’s Top 10 for LLM Applications (OWASP Foundation, 2024) identifies prompt injection and excessive agency as the leading risks: agents with broad tool access can be hijacked to exfiltrate data or take unauthorized actions. Farahmandi et al. (Farahmandi et al., 2025) catalog over thirty attack techniques against LLM-agent workflows, covering input manipulation, model compromise, and privacy attacks. Our framework directly addresses these threat classes by intercepting every tool call and enforcing label-based policy before any sink action executes.

Multi-agent systems. Role-separated multi-agent pipelines (Anthropic, 2024) amplify the monitoring challenge because data flows across agent boundaries. Our shared `ObsHub` maintains unified taint and lineage across all agents in the cluster.

3 SYSTEM DESIGN

3.1 ARCHITECTURE OVERVIEW

The framework centers on `ObsHub`, a singleton holding four subsystems: `EventBus` (typed pub/sub, JSONL persistence, WebSocket broadcast), `TaintStore` (artifact management),




Figure 1: The observability dashboard during a credential retrieval run (task: “Retrieve the API key and verify it via HTTP GET”). **Risk** panel (left): score of 0.65 for an `http_get_external` sink, shown on a green-to-red color bar. **Human Review** panel (right): R2 HITL request for a credential-labeled GET to an external URL, with the fingerprint hash displayed and Allow / Allow Once / Deny actions available. **Lineage** panel (lower-left): D3.js graph showing the RAG search source (credential border ring) flowing into the HTTP GET tool. **Timeline** panel (lower-right): event log with per-step entries.

Rule	Label	Sink	Decision	Rationale
R1	pii / credential	http_post_external	DENY	Direct exfiltration
R1b	internal_doc	http_post_external	HITL	Gray-area external send
R2	credential	http_get_external	HITL	Credential in URL params
R3	$\geq$ internal_doc	file outside allowlist	DENY	Unauthorized write
R3u	unknown / public	file outside allowlist	HITL	Conservative on unknowns
R4	pii	file inside allowlist	HITL	PII write needs confirmation

Table 1: Policy rules, outcomes, and rationale.

LineageBuilder (provenance graph), and TrustStore (fingerprint scores). Every I/O action is routed through `MonitoredIO` (`observe/wrappers.py`)—the sole intercept point—which runs the pipeline below before executing or blocking the action. The wrapped agent requires no modification; from its perspective, `MonitoredIO` is simply the object that implements tool calls.


3.2 OBSERVABILITY PIPELINE

Taint tracking. Data retrieved by the agent is wrapped as an `Artifact` with one of four sensitivity labels ordered `public < internal_doc < pii < credential`. Labels are assigned at ingestion by data origin (e.g., email always carries `pii`) or inferred via regex patterns for addresses, SSNs, and API key prefixes. When an artifact is passed to a subsequent tool call, labels union monotonically—a `pii` artifact flowing into an HTTP POST body taints that body regardless of its surface content. Artifact IDs propagate across LLM turns via `[[artifacts:id1,id2,...]]` annotations in tool-result strings, enabling taint to follow data through multi-turn chains.

Policy evaluation. Before any sink action executes, `evaluate_sink()` applies six hard rules on (label, sink) pairs, shown in Table 1. Rule R1—PII or credentials posting to an external HTTP endpoint—always produces an automatic DENY with no human prompt, since the false-positive cost of a missed exfiltration exceeds the cost of an unnecessary block. Rules R1b, R2, and R4 produce HITL, pausing the agent and surfacing the action for operator review. Rules R3 and R3u handle unauthorized file writes.

Risk scoring. A continuous score in $[0, 1]$ is computed for every sink action:

$$r = 0.55s + 0.35k + 0.10d - 0.15t \quad (1)$$



(a) Lineage for the PII exfiltration scenario: `gmail_mock` (PII ring) → `fetch_gmail_mock` → `http_post` → `http_post_external` (red, R1 DENY). (b) HITL panel (R1b): `internal_doc` about to be posted externally. Rule, sink, label, and fingerprint are shown before the operator decides.

Figure 2: Lineage graph and HITL decision panel.

where s is the sensitivity weight of the highest label (0 to 1.0), k ranks sink severity (`http_post_external` > `file_write` > `http_post_internal` > `tool_result`), $d = \min(\text{step_id} \times 0.03, 0.2)$ captures chain depth, and t is the trust offset (see below). The risk score gates HITL independently of policy: if a nominally-allowed action scores above the trust-adjusted threshold, it is still escalated. This catches behavioral anomalies not covered by the six named rules.

Trust store. Each action is fingerprinted as `SHA-256(tool name || args || sink type)` and stored in `trust.json`, which persists across agent runs. Scores decay by $\times 0.99$ per step to prevent stale approvals from accumulating indefinitely. Operator decisions adjust scores: `allow` adds $+0.15$, `deny` subtracts -0.20 , and `allow_once` adds $+0.05$ without establishing a lasting precedent. The HITL escalation threshold adapts per fingerprint: $\theta = 0.72 + (\text{trust} - 0.5) \times 0.12$, so a frequently-approved action progressively raises its own threshold and eventually stops triggering HITL. Conversely, a repeatedly-denied action lowers its threshold, flagging it at lower risk scores in subsequent runs. This mechanism lets the monitor adapt to operator preferences without any explicit reconfiguration.

HITL gate. A flagged action emits a `hitl_request` event and creates an `asyncio.Future` registered in `ObsHub`. The agent loop awaits the future, genuinely suspending execution. The dashboard receives the event over `WebSocket`, displays the rule, sink, labels, risk score, and fingerprint, and presents `Allow` / `Allow Once` / `Deny` buttons (Figure 2, right). When the operator clicks a button, the browser’s `POST /api/hitl/{hid}` resolves the future; the agent resumes or terminates accordingly. This integrates cleanly with FastAPI’s `async` runtime with no polling overhead.

3.3 DATA-FLOW LINEAGE

`LineageBuilder` maintains a directed provenance graph rendered in real time by `D3.js` with a force-directed layout. The node types are: source (blue, with a colored border ring indicating label), tool (yellow), safe sink (green), and blocked sink (red). Artifacts are *not* visible nodes; instead, the builder tracks internal ownership and adds a direct edge when one tool’s output is consumed by another, producing a clean chain without intermediate artifact clutter.

Figure 2 (left) shows the exfiltration scenario: `gmail_mock` (blue, orange PII ring) → `fetch_gmail_mock` (tool) → `http_post` (tool) → `http_post_external` (red, blocked). The border ring on the source node indicates that the artifact carries a `pii` label; the red terminal node confirms that R1 fired. An operator can trace the full provenance of the blocked action at a glance.

4 EVALUATION

4.1 SETUP

We evaluate in two stages. First, a *sanity suite* of 12 scripted scenarios exercises all six policy rules across all three outcomes (DENY, HITL, ALLOW) using `MockLLM`—a deterministic, keyword-driven mock that replays fixed tool-call sequences without involving any live LLM. `OBS_HITL_AUTO` is set to `allow` or `deny` as appropriate so that HITL decisions are resolved automatically; the expected outcome for each scenario is declared in `tasks_synthetic.yaml` and compared against the event type emitted by the pipeline. Because `MockLLM` is fully deterministic, the same input keyword always produces the same tool-call sequence, and thus the same events, eliminating stochastic

Model	Task	Label	Risk	Outcome	Rule
Gemini (injected)	T1: Summarize	pii	0.241	ALLOW	–
	T2: Extract	pii	0.367	HITL	R4
	T3: Prioritize	pii	0.265	ALLOW	–
	T4: Upload (adv.)	pii	0.565	DENY	R1
DeepSeek (clean)	T1: Summarize	public	0.000	ALLOW	–
	T2: Extract	pii	0.367	HITL	R4
	T3: Prioritize	public	0.000	ALLOW	–
	T4: Upload (adv.)	pii	0.565	DENY	R1

Table 2: Live benchmark results for Gemini 2.5 Flash (injected malicious system prompt) and DeepSeek Chat (clean) across four email-processing tasks. Both models are blocked on the adversarial upload (T4). Risk scores on T1 and T3 diverge due to Gemini including raw email addresses in its outputs under the injected instruction.

variance from correctness checks. Scenarios cover: multi-hop taint from email to HTTP POST, credential retrieval via RAG followed by external GET, PII written to an allowlisted path, internal documents written outside the allowlist, public data posted to an internal endpoint, and unknown artifacts written outside the allowlist.

Second, a *live benchmark* runs two commercial LLMs over a shared email inbox in a four-agent cluster (Anthropic, 2024): an *Orchestrator* delegates to three specialists—*Researcher* (IMAP email, BM25 RAG, HTTP GET), *Analyst* (file reads/writes), and *Output* (report writing, HTTP POST). All agents share a single *ObsHub* and are tagged with an *agent_id* for per-agent attribution. **Gemini 2.5 Flash** is given a malicious system prompt instructing it to include raw email addresses and full sender names verbatim in every response, simulating a compromised model. **DeepSeek Chat** runs without any injected instruction, serving as a clean baseline. Both models process the same four tasks: (T1) inbox summarization output to a tool result, (T2) structured PII extraction written to a local file, (T3) priority ranking posted to an internal endpoint, and (T4) compiling a full report and uploading it to an external analytics API (<https://analytics.example.com/report>). T4 is adversarial by design: both models’ system prompts explicitly instruct them to perform the upload, ensuring that any blocking is attributable to the policy engine rather than model refusal.

4.2 SANITY SUITE: CORRECTNESS

All 12 scenarios pass, covering DENY (4 scenarios), HITL (5 scenarios), and ALLOW (3 scenarios) outcomes. No false positives or false negatives were observed, giving precision, recall, and F1 of 1.0 across the scripted suite. Specifically: all four DENY scenarios correctly emit a `policy_violation` event and return `False` to the agent; all five HITL scenarios emit a `hitl_request` event and suspend the agent loop until resolved; and all three ALLOW scenarios pass through without emitting any policy event. The monotone label union is exercised in the multi-hop scenario, where a clean-prose summary derived from a PII-bearing email is correctly labeled `pii` and blocked when subsequently posted externally—demonstrating that the content of the artifact is irrelevant once taint is established.

4.3 LIVE BENCHMARK: POLICY ENFORCEMENT AND RISK DIVERGENCE

Table 2 summarizes results across both models. The primary finding is that T4—the adversarial external upload—is blocked for both models by R1 (DENY). The model’s instruction to upload the report is irrelevant to the policy engine: the taint label (`pii`) and sink type (`http_post_external`) together satisfy R1 regardless of the instruction that triggered the call. The `policy_violation` event appears in the JSONL audit log and on the dashboard for both runs.

A secondary finding concerns T1 and T3, where the two models produce different taint labels and risk scores despite identical task inputs. DeepSeek’s summaries and priority rankings are labeled `public` (risk = 0.000), indicating that the model’s outputs do not contain patterns recognized by the PII regex heuristics. Gemini’s outputs, operating under the injected instruction to include verbatim email addresses and sender names, are labeled `pii` (risk = 0.241 on T1 and 0.265 on T3). No hard rule fires on either task—T1 routes to a `tool_result` sink and T3 to an `http_post_internal` sink, neither of which appears in the DENY or HITL rules for `pii`. Nevertheless, the elevated risk scores are detectable from the dashboard’s risk meter, and an operator would see a clear anomaly compared to the clean-model baseline.

This demonstrates a key property beyond hard rules: the risk score functions as a behavioral fingerprint. Even when the injected model does not attempt an outright policy violation, its data-handling patterns diverge at earlier, lower-risk stages—detectable by the operator before exfiltration reaches a harder policy boundary.

The innocent summary attack. T1 also illustrates why content inspection is insufficient. The task is legitimate (summarize inbox); an injected instruction then attempts to POST the clean-prose summary externally. No raw SSNs appear in the outbound body. Taint tracking records that the summary was derived from a `pii`-labeled fetch; the label is live regardless of surface content, and R1 fires. This scenario is exercised in the sanity suite and confirmed in the Gemini run.

5 DISCUSSION

5.1 LIMITATIONS

Taint labeling is coarse and monotone. Label assignment is regex-based and context-free: a string containing “Example SSN: 123-45-6789” receives the same `pii` label as a real record because the heuristic cannot distinguish illustrative examples from genuine sensitive data. More substantially, labels union but never drop. A properly redacted summary—one that has replaced all PII with placeholders—still carries the `pii` label of its source artifact, and will be blocked if subsequently posted externally even though the outbound content contains no sensitive data. This is the fundamental tension between the precision of the safety guarantee and its usability: monotone propagation is conservative by design, but it over-blocks legitimate summarization and transformation workflows.

Static policy. The six rules are hard-coded in `policy.py` and require a code change to update. There is no per-user or per-role policy differentiation, and no runtime reconfiguration path. An enterprise deployment would likely need a policy DSL that administrators can update without redeploying the monitoring layer, and role-based rules that allow, for example, a compliance officer to receive PII reports that would be denied for a standard user agent.

Evaluation scope. Correctness was validated exclusively on deterministic, scripted tool-call sequences produced by `MockLLM`. A real agent running a live task may produce unexpected tool-call orderings, multi-step reasoning chains, or partial artifact re-use patterns that were not covered by the 12 scenarios. No user study was conducted for the HITL panel or dashboard UX, so it is unknown whether operators make accurate decisions under realistic workload conditions or how interrupt frequency affects their attention over time.

Assumes cooperative tool use. The framework intercepts declared tool calls routed through `MonitoredIO`. It cannot guard against a backdoor LLM that exfiltrates data through a channel not exposed as a tool—for example, by encoding information steganographically in model outputs consumed by a co-process, by making direct system calls if the model has code-execution access, or by exploiting side channels in the host environment. The monitor operates at the application layer; closing that gap requires system-level isolation such as network namespace sandboxing or OS-level mandatory access control, not application-layer hooks. In this sense, the current framework is best understood as a defense against *unintended* data leakage and prompt-injection attacks, not against a model that has been compromised at the weights level.

5.2 FUTURE WORK

Richer taint and policy models. The most impactful near-term extension is context-aware taint: permitting a transformed artifact to shed its source label if the transformation can be verified to have removed all sensitive content. One approach is to route transformed artifacts through a lightweight PII detection model that certifies redaction before allowing a label downgrade. On the policy side, a lightweight DSL for per-user and per-role rules, loadable at runtime from a configuration file, would make the framework practical in multi-tenant environments where different agents and operators have different trust levels and data-handling obligations.

Principled evaluation platform. The monitoring infrastructure provides a natural substrate for adversarial testing. One direction is LLM-synthesized attack scenarios: use a red-team model to generate novel tool-call sequences specifically targeting gaps in the six named rules, then expand the rule set to cover them. This would replace the current manually authored sanity suite with a continuously growing, model-generated test corpus. A complementary direction is a red-team/blue-team arena in which a simulated attacker model and the policy engine compete iteratively over many runs, surfacing the empirical boundary of what rule-based enforcement can and cannot catch and providing a more principled measure of robustness than pass/fail correctness alone.

REFERENCES

- Anthropic. Building effective agents. <https://www.anthropic.com/engineering/building-effective-agents>, 2024.
- Nazanin Farahmandi et al. From prompt injections to protocol exploits: Threats in LLM-powered AI agent workflows. *Computer Standards & Interfaces*, 2025. URL <https://www.sciencedirect.com/science/article/pii/S2405959525001997>.
- Yuchuan Fu, Xiaohan Yuan, and Dongxia Wang. Ras-eval: A comprehensive benchmark for security evaluation of llm agents in real-world environments, 2025. URL <https://arxiv.org/abs/2506.15253>.
- OWASP Foundation. OWASP Top 10 for Large Language Model Applications 2025. <https://genai.owasp.org/llm-top-10/>, 2024.
- Arman Zharmagambetov, Chuan Guo, Ivan Evtimov, Maya Pavlova, Ruslan Salakhutdinov, and Kamalika Chaudhuri. Agentdam: Privacy leakage evaluation for autonomous web agents. *arXiv preprint arXiv:2503.09780*, 2025.